

# Advanced Topics in Malware Analysis

## Execution Tracing

**Brendan Saltaformaggio, PhD**

*Assistant Professor*

School of Electrical and Computer Engineering

Tracing: Source Level Instrumentation



1

## Lesson Objectives

- Employ the execution tracing for dynamic analysis
- Apply static and dynamic binary instrumentation
- Create a Pintool to utilize PIN
- Discuss Valgrind and QEMU techniques for instrumentation



2

## The Dynamic Analysis Swiss Army Knife

- Tracing is a process that faithfully records detailed information of program execution
  - **Benefits:** Lossless, simple to implement, low overhead
  - **Problems:** Requires heavy after-the-fact analysis on the execution trace
- Control flow tracing
  - Logs the sequence of **executed** statements
- Dependence tracing
  - Logs the sequence of **exercised** dependencies
- Value tracing
  - Logs the sequence of **values** that are produced by each instruction
- Memory access tracing
  - Logs the sequence of **memory** references during an execution



3

## You Have Already Performed Tracing!

Tracing by **printf**

```
int max = 0;
for (p = head; p; p = p->next)
{
    printf("In the loop!\n");
    if (p->value >= max)
    {
        printf("True branch.\n");
        max = p->value;
    }
}
```

This is the essence of dynamic tracing!

May seem silly, but...

Consider this output:

```
In the loop!
In the loop!
True branch.
In the loop!
True branch.
In the loop!
True branch.
```

- How many elements in the list?
- How many negative values?



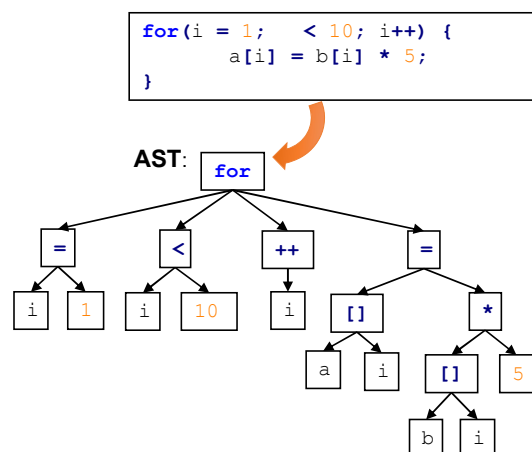
4

## Tracing by Source Level Instrumentation

- Typically performed via a compiler plugin
  - LLVM, GCC plugin, etc.
- 1. Read a source file and parse it into **abstract syntax trees (ASTs)**
- 2. Annotate the ASTs with instrumentation
- 3. Translate the instrumented ASTs into a new source file
- 4. Compile the new source
- 5. Execute the instrumented program and a trace is produced

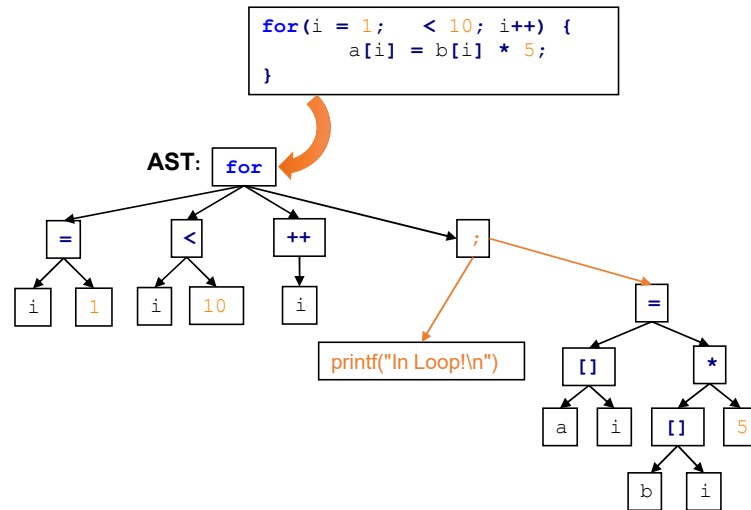
5

## Tracing by Source Level Instrumentation



6

## Tracing by Source Level Instrumentation



7

## Limitations of Source Level Instrumentation

- **Requires source code**
  - Worms and viruses are rarely provided with source code
  - Closed source software
  - That is why we are in this class ☺
- **Hard to handle libraries**
  - Proprietary libraries: communication (MPI, PVM), linear algebra (NGA), database query (SQL libraries)
- **Hard to handle multi-lingual programs**
  - Source code level instrumentation is language dependent
  - Have to rewrite everything for Java, C++, Python, ...

8

# Advanced Topics in Malware Analysis

## Execution Tracing

**Brendan Saltaformaggio, PhD**

*Assistant Professor*

School of Electrical and Computer Engineering

Tracing: Binary Instrumentation



9

## Tracing by Binary Instrumentation

- The **most important** binary-analysis capability to date: **Binary Instrumentation**
- **Two flavors:**
  1. Static Binary Instrumentation
  2. Dynamic Instrumentation
- **Features:**
  - No source code required
  - Directly handles library binaries (either statically instrumenting or during their execution)
  - Does not care what language(s) the program is written in --- any binary can be instrumented
    - **Worst case example:** Instrument the JVM & monitor the Java program during execution



10

## Static Binary Instrumentation

- Insert instrumentation into the binary executable
  - Instrumentation will run next time the binary is executed
  - A.K.A. Binary “Rewriting”
1. Given a binary executable, **parse** it into **intermediate** representation
    - Could be as simple as instruction models (like we did in Lab 3 & Lab 4)
    - More advanced representations such as control flow graphs may also be generated
  2. Design tracing instrumentation for the intermediate representation
    - E.g., Control flow tracing → for each (instruction in binary): `printf(“%i”, instruction.address);`
  3. A lightweight “compiler” inserts the instrumentation logic into a new executable



11


## Static Binary Instrumentation Example

Original Source:

```
#include <stdio.h>

int main(int argc, char* argv[]) {
    if (argc == 2)
        printf("Hello %s\n", argv[1]);
    return 0;
}
```

Original Binary  
(64-bit Linux with -O3)



```
1.      cmp     edi, 2
2.      jz      print
3.      xor     eax, eax
4.      retn
5. print: push     rax
6.      mov     rsi, [rsi+8]
7.      mov     edi, offset format ; "Hello %s\n"
8.      xor     eax, eax
9.      call    _printf
10.     xor     eax, eax
11.     pop     rdx
12.     retn
```

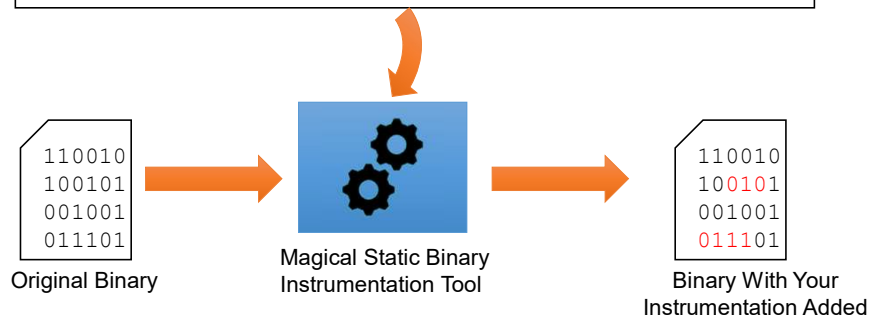


12

## Static Binary Instrumentation Example

Your Instrumentation Logic:

```
void trace_ctrl_flow(inst *current_inst) {
    printf("%u\n", current_inst->address);
}
...
for(inst *in = get_instructions(); in; in = in->next) {
    add_instrumentation(in, trace_ctrl_flow);
}
```



13

## Static Binary Instrumentation Example

```
inst:    push    rsi    ; save all dirty regs
         push    rax
         push    rflags
         mov     rsi, rdi ; embedded instrumentation
         mov     edi, offset u ; "%u\n"
         call    _printf
         pop     rflags ; restore dirty regs
         pop     rax
         pop     rsi
         ret

...
1.       cmp     edi, 2
         push    rdi    ; call instrumentation
         mov     rdi, 1
         call    inst
         pop     rdi
         push    rdi    ; call instrumentation
         mov     rdi, 2
         call    inst
         pop     rdi
2.       jz      print  ; be careful on ctrl flow!
```

14

## Static Binary Instrumentation Example

```

3.      xor     eax, eax
        push    rdi    ; call instrumentation
        mov     rdi, 3
        call    inst
        pop     rdi
        push    rdi    ; call instrumentation
        mov     rdi, 4
        call    inst
        pop     rdi
4.      retn
5. print: push    rax
        push    rdi    ; call instrumentation
        mov     rdi, 5
        call    inst
        pop     rdi
6.      mov     rsi, [rsi+8]
        push    rdi    ; call instrumentation
        mov     rdi, 6
        call    inst
        pop     rdi

```



15

## The Good & Bad of Static Binary Instrumentation

- Almost impossible to do accurately!
  - Pointers must be carefully preserved or adjusted
  - Memory data must not be corrupted
  - Original disassembly must be 100% correct
    - You cannot instrument data that was misinterpreted as code!
- Suggested readings:
  - Deng, Z., Zhang, X., & Xu, D. (2013). BISTRO: Binary Component Extraction and Embedding for Software Security Applications. *Proceedings of the 18th European Symposium on Research in Computer Security- ESORICS 2013*.
  - Wartell, R., Mohan, V., Hamlen, K. W., & Lin, Z. (2012). Securing untrusted code via compiler-agnostic binary rewriting. *Proceedings of the 28th Annual Computer Security Applications Conference- ACSAC 12*.
- Significantly increases executable file size
- Good: FAST! All the instrumentation logic is “baked in” to the executable



16



## Binary Instrumentation: Static vs Dynamic

- **Static Binary Instrumentation:** Given an original binary executable and generate an instrumented executable that can be executed with our analysis embedded within!
  - Instrument statically = before runtime
- **Dynamic Binary Instrumentation:** Given an original binary executable and an input, start executing the binary with the input, and during execution add instrumentation to the binary on the fly
  - Instrument dynamically = during runtime
- **Advantages for Dynamic Instrumentation:**
  - No need to recompile or relink
  - Discover code at runtime
  - Handle dynamically-generated code
  - Attach to running processes



17


## Dynamic Binary Instrumentation Example

- Original Source:

```
#include <stdio.h>

int main(int argc, char* argv[]) {
    if (argc == 2)
        printf("Hello %s\n", argv[1]);
    return 0;
}
```

Original Binary  
(64-bit Linux with -O3)



```
1.      cmp     edi, 2
2.      jz      print
3.      xor     eax, eax
4.      retn
5. print: push     rax
6.      mov     rsi, [rsi+8]
7.      mov     edi, offset format ; "Hello %s\n"
8.      xor     eax, eax
9.      call    _printf
10.     xor     eax, eax
11.     pop     rdx
12.     retn
```

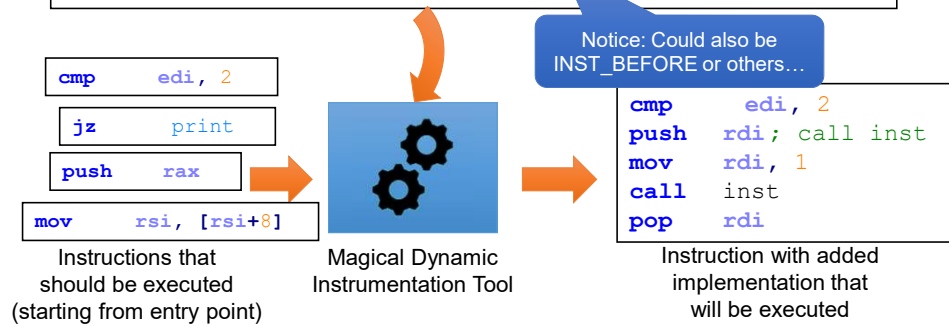


18

## Dynamic Binary Instrumentation Example

Your Instrumentation Logic:

```
void trace_ctrl_flow(inst *current_inst) {
    printf("%u\n", current_inst->address);
}
...
for(inst *in = get_instructions(); in; in = in->next) {
    add_instrumentation(in, INST_AFTER, trace_ctrl_flow);
}
```



Georgia Tech

19

## Dynamic Binary Instrumentation

- A few flavors.
- We will look at:
  - **Platform-specific** --- instrumentation targets the exact instructions that are executing on the CPU
  - **Emulation-based** --- instrumentation targets some intermediate language
- Others (some are hybrids of these categories)

DynamoRIO  
The DR. is in.

Dyninst

QEMU

Pin

Valgrind

Georgia Tech

20

# Advanced Topics in Malware Analysis

## Execution Tracing

**Brendan Saltaformaggio, PhD**

*Assistant Professor*

School of Electrical and Computer Engineering

PIN



21

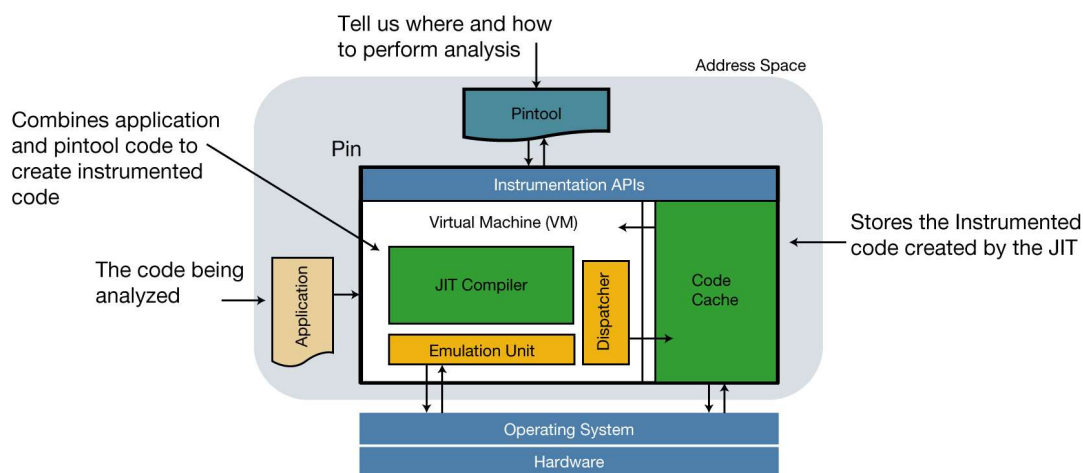
## The King of x86 Dynamic Instrumentation: PIN

- **Pin** is a platform for dynamic instrumentation of binary programs
- **Pin** is maintained by **Intel**
- Supports Intel architectures: **x86, x86-64, Itanium, Xscale**
- Can instrument executables on a variety of OS's: Windows, Linux, OSX, Android
- You write a "**Pintool**" which instructs Pin on how you want to instrument the execution
- Pintools are very powerful (that is why Pin is king) because they can **insert arbitrary code** (written in C or C++) at **arbitrary locations** in an executable
- Instrumentation code is added dynamically while the executable is running
  - It can also be changed on-the-fly!
  - Handles threads & asynchronous signals!
  - This makes it possible to attach Pin to an already running process



22

## The King of x86 Dynamic Instrumentation: PIN



Read More: Pin: building customized program analysis tools with dynamic instrumentation. In ACM SIGPLAN Notices (Vol. 40, No. 6, pp. 190-200).



23

## PinTools

- <https://software.intel.com/en-us/articles/pin-a-dynamic-binary-instrumentation-tool>
- Launch and instrument an application:

`$ pin -t pintool.so --application`

Instrumentation engine (provided in the kit)      Instrumentation tool(write your own, or use one provided in the kit)

- Attach to and instrument an application: `$ pin -t pintool.so -pid 1234`
- Like IDA, Pintools are specific to 32-bit and 64-bit!
  - May share code (watch out for incompatibilities!), but must be compiled separately



24

## Pintools API

- Basic APIs are architecture independent
- Provide common functionalities like determining:
  - Control-flow changes
  - Memory accesses
- Architecture-specific APIs
  - E.g., info about opcodes and operands
- State change APIs
  - E.g., process create/end, interrupt received
- Call-based APIs:
  - Instrumentation routines
  - Analysis routines



25

## Pintools Structure

Two types of core routines:

- 1. Instrumentation routines** define where instrumentation is inserted
  - E.g., before instruction
  - Occurs first time an instruction is executed
  - Other APIs can remove instrumentation after it has been inserted
- 2. Analysis routines** define what to do when instrumentation is activated
  - E.g., print instruction address, increment counter
  - Occurs every time an instruction is executed



26

## Pintool Example: ManualExamples/itrace.cpp



27

```
#include <stdio.h>
#include "pin.h"
FILE * trace;

void printip(void *ip) {
    fprintf(trace, "%p\n", ip);
}

void Instruction(INS ins, void *v) {
    INS_InsertCall(ins, IPOINT_BEFORE,
        (AFUNPTR)printip, IARG_INST_PTR,
        IARG_END);
}

void Fini(INT32 code, void *v) {
    fclose(trace);
}

int main(int argc, char * argv[]) {
    trace = fopen("itrace.out", "w");
    PIN_Init(argc, argv);
    INS_AddInstrumentFunction(Instruction, 0);
    PIN_AddFiniFunction(Fini, 0);
    PIN_StartProgram();
    return 0;
}
```

## Pintool Example: ManualExamples/itrace.cpp

Analysis  
Routine

Instrumentation  
Routine

State change  
Analysis Routine

Main calls the Pin  
APIs to register the  
instrumentation  
routines



28

## Why PIN is King

- Versatility! Pin can do sooooo many things
- Sample tools in the Pin distribution:
  - See **C:\pin\source\tools** directory!!
  - Cache simulators, branch predictors, address tracer, syscall tracer, edge profiler, stride profiler
- Some tools developed and used inside Intel:
  - **Opcodemix** --- analyze code generated by compilers
  - PinPoints --- find representative regions in programs to simulate
- Companies are writing their own Pintools
- Universities use Pin in teaching and research
  - My PhD was one big journey in Pin --- <https://github.com/bdsaltaformaggio/DSCRETE>
  - Many others...



29

## Advanced Topics in Malware Analysis

### Execution Tracing

### Brendan Saltaformaggio, PhD

*Assistant Professor*

School of Electrical and Computer Engineering

Valgrind



30

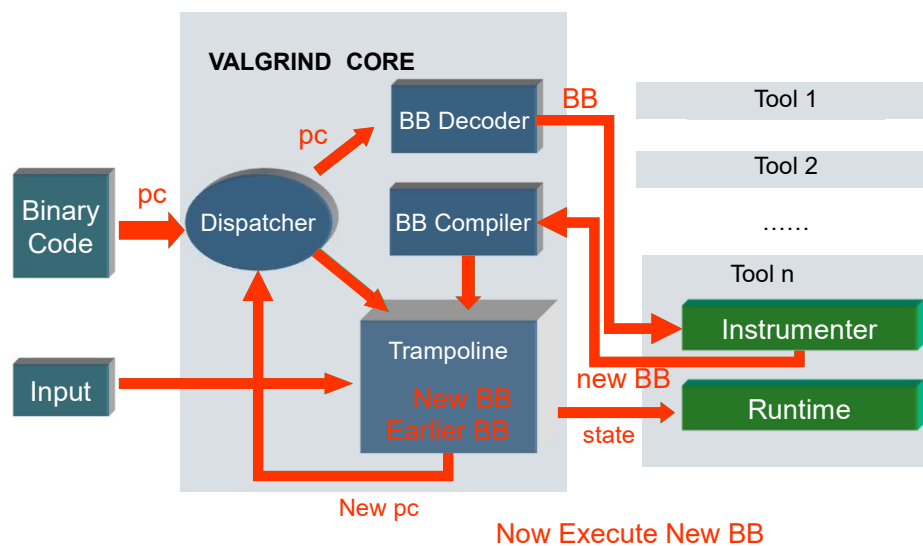
## Dynamic Binary Instrumentation - Valgrind

- Developed by Julian Seward at Cambridge University
  - Google - O'Reilly Open Source Award for "Best Toolmaker" 2006
  - A merit (bronze) Open Source Award 2004
  - Open source (Pin is not!)
  - Works on many platforms: x86, PowerPC, ARM, ...
- Overhead is a big problem
- 5-10 **\*times\*** slowdown **without** any instrumentation
  - Pin has very little slowdown for no instrumentation
- Suggested reading:
  - Nethercote, N., & Seward, J. (2007). Valgrind: A Framework for Heavyweight Dynamic Binary Instrumentation. *Proceedings of the 2007 ACM SIGPLAN Conference on Programming Language Design and Implementation - PLDI 07*.



31

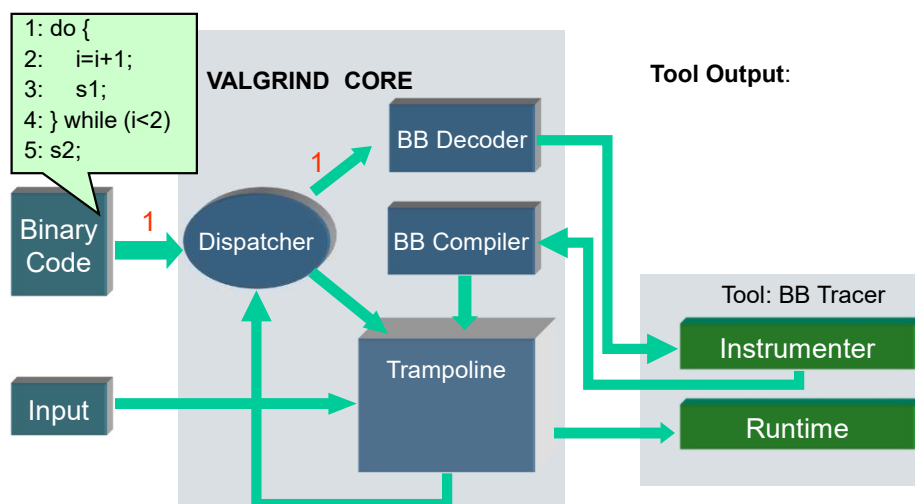
## Valgrind Infrastructure



32



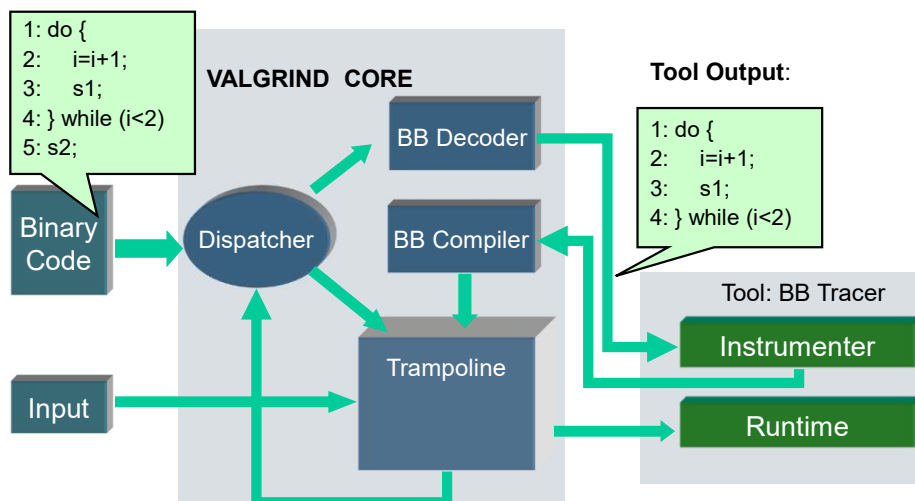
## Valgrind Example



Georgia  
Tech

33

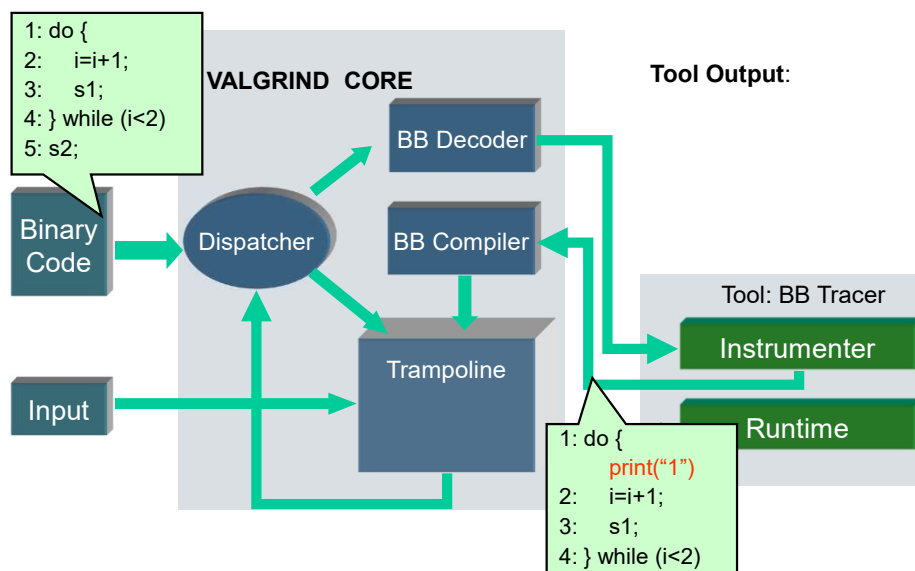
## Valgrind Example



Georgia  
Tech

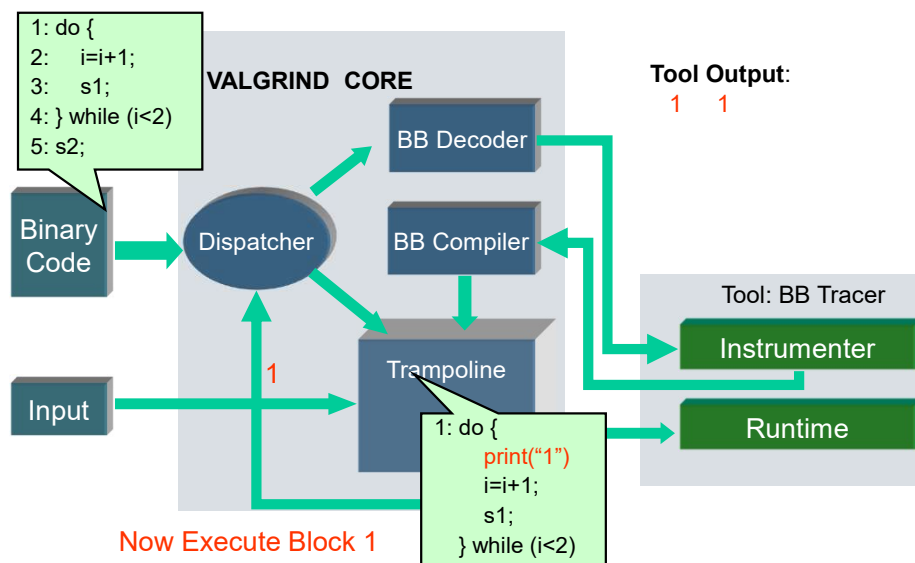
34

## Valgrind Example



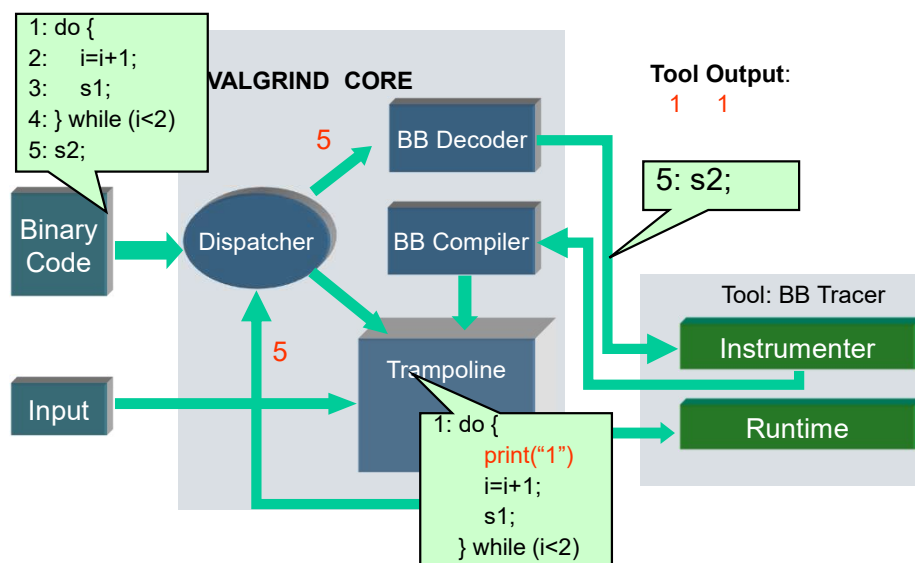
35

## Valgrind Example



36

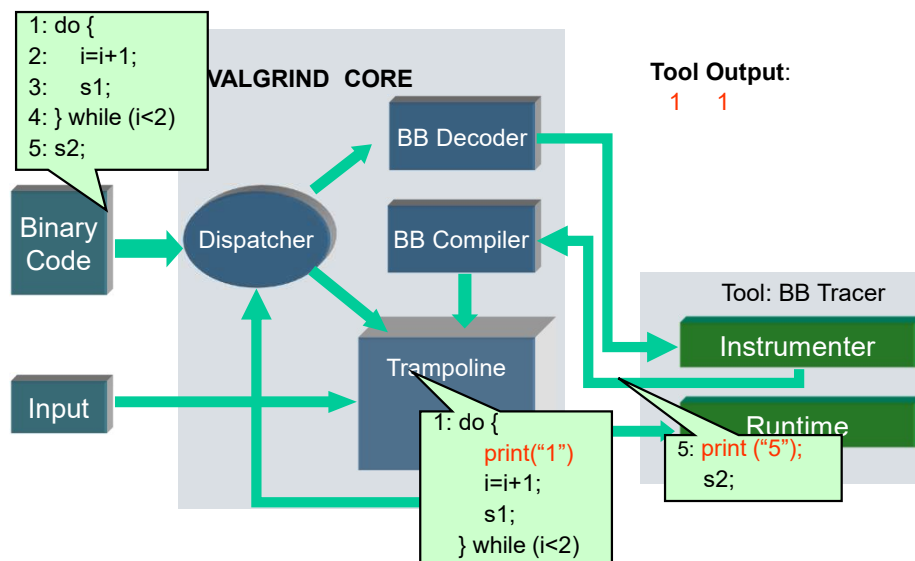
## Valgrind Example



Georgia Tech

37

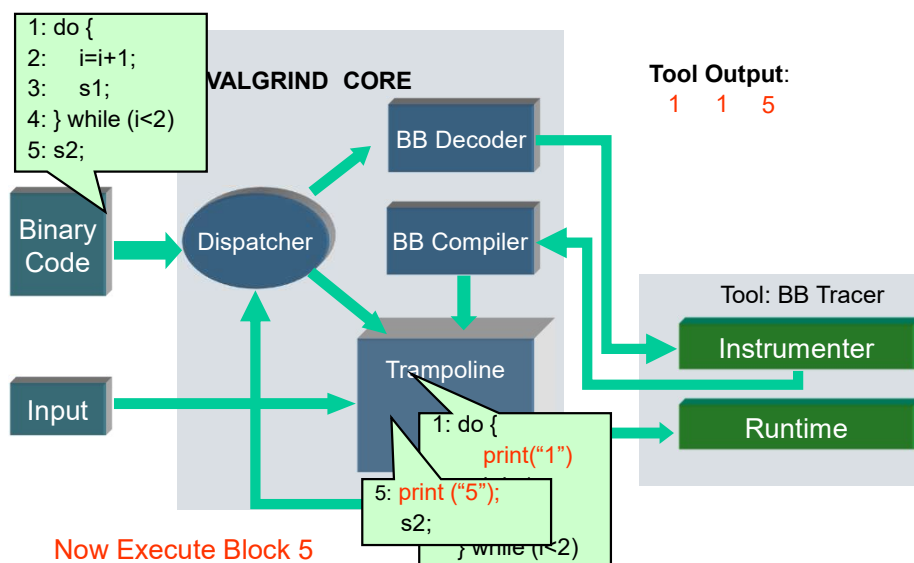
## Valgrind Example



Georgia Tech

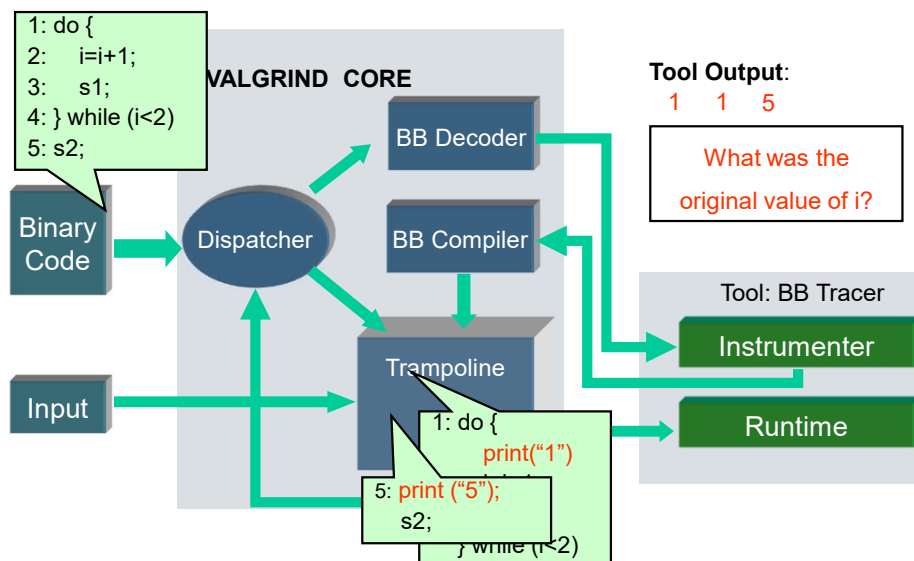
38

## Valgrind Example



39

## Valgrind Example



40

# Advanced Topics in Malware Analysis

## Execution Tracing

**Brendan Saltaformaggio, PhD**

*Assistant Professor*

School of Electrical and Computer Engineering

QEMU



41

## QEMU Emulation-based Instrumentation

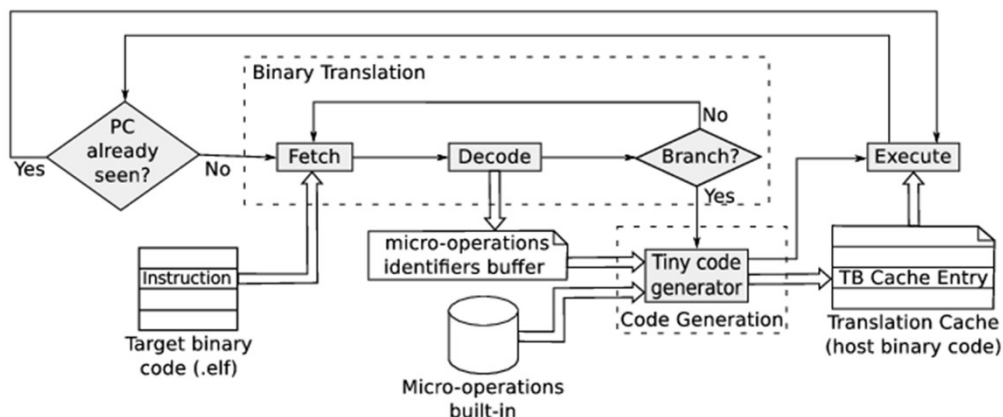
- QEMU is a generic and open source machine **emulator** and hypervisor (VM manager)
- As a **machine emulator**, QEMU can run OSes and programs made for one machine (e.g. an ARM board) on a different machine (e.g. your own PC)
  - By using dynamic translation, it achieves very good performance
- As a **hypervisor**, QEMU achieves near native performance by executing the guest code directly on the host CPU
  - QEMU supports virtualization when executing under the Xen hypervisor or using the KVM kernel module in Linux
  - When using KVM, QEMU can virtualize x86, server and embedded PowerPC, and S390 guests
- In either case, you can add your analysis routines to QEMU's code base to instrument the entire guest OS as it is executing!



42

## QEMU Infrastructure

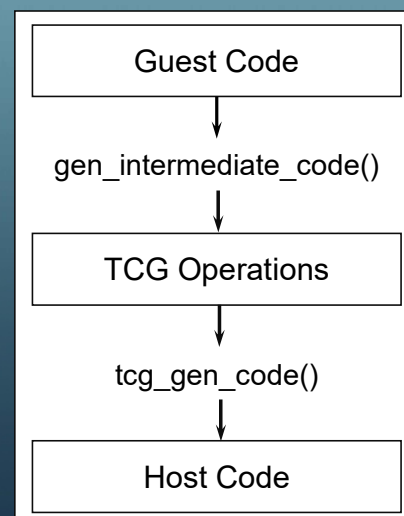
- How does QEMU execute ARM code on x86?



43

## QEMU Code Translation

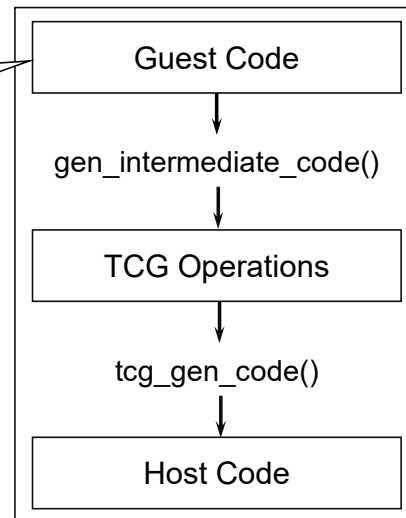
- QEMU uses an intermediate representation (IR)
- Frontends are in the "target" directory
  - These are guest architectures
  - "What QEMU executes"
  - Includes alpha, arm, cris, i386, m68k, mips, ppc, sparc, ...
- Backends are in tcg/
  - "Tiny Code Generator"
  - These are host CPU operations
  - "Where QEMU executes"
  - Includes arm/, i386/, ia64/, mips/, ppc/, ppc64/, s390/, sparc/, ...



44

## QEMU Code Translation Example x86 Guest -> x86 Host

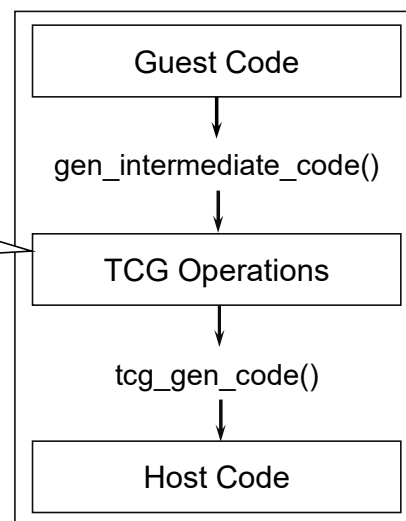
```
push    ebp
mov     ebp, esp
not     eax
add     edx, eax
mov     eax, edx
xor     eax, 0x55555555
pop     ebp
ret
```



45

## QEMU Code Translation Example x86 Guest -> x86 Host

```
ld_i32      tmp2,env,$0x10
qemu_ld32u  tmp0,tmp2,$0xffffffff
ld_i32      tmp4,env,$0x10
movi_i32    tmp14,$0x4
add_i32     tmp4,tmp4,tmp14
st_i32      tmp4,env,$0x10
st_i32      tmp0,env,$0x20
movi_i32    cc_op,$0x18
exit_tb     $0x0
```



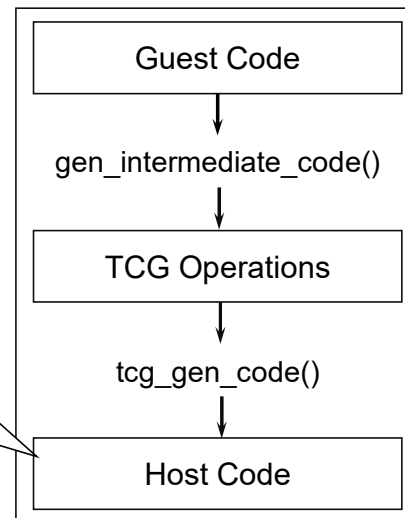
46

## QEMU Code Translation Example x86 Guest → x86 Host

```

mov 0x10(%ebp), %eax
mov 0x10(%ebp), %edx
mov (%ecx), %eax
mov %eax, %ecx
add $0x4, %edx
mov %edx, 0x10(%ebp)
mov %eax, 0x20(%ebp)
mov $0x18, %eax
mov %eax, 0x30(%ebp)
xor %eax, %eax
jmp 0xba0db428
. . .
/* This represents just
the ret instruction! */


```



47

## Qemu Block Chaining

- Like we saw in Valgrind, returning to the dispatcher from the code cache is very slow!
- Solution: jump directly between basic blocks!
  - Make space for a jump, followed by a return to the epilogue
  - Every time a block returns, try to chain it



48



# Advanced Topics in Malware Analysis

## Execution Tracing

**Brendan Saltaformaggio, PhD**

*Assistant Professor*

School of Electrical and Computer Engineering

Offline Trace Reconstruction



49

## Key Word: Lossless

- Dynamic tracing is assumed to be lossless
  - NO processing of output during runtime!
  - All analysis & processing takes place offline
- A trace must allow offline analysis tool to faithfully recreate the analysis target
  - **Control flow tracing** --- sequence of executed statements
  - **Dependence tracing** --- sequence of exercised dependencies
  - **Value tracing** --- sequence of values that are produced by each instruction
  - **Memory access tracing** --- sequence of memory references during an execution
- Therefore, tracing is the most fine-grained of all possible dynamic analyses



50

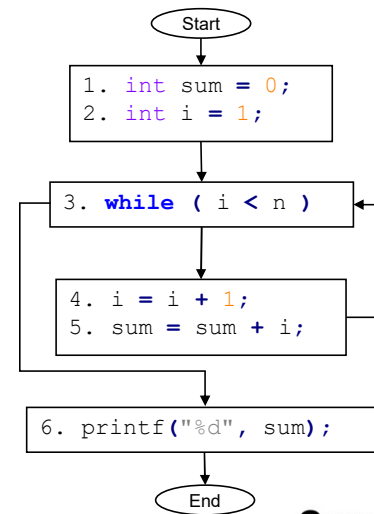
## Fine-Grained Tracing is Expensive

```
void sumUp(int n) {
    int sum = 0;
    int i = 1;
    while ( i < n ) {
        i = i + 1;
        sum = sum + i;
    }
    printf("%d",
sum);
}
```

**Trace(n=6) = ?**

**Trace(n=6):** 1 2 3 4 5 3 4 5 3 4 5 3 4 5 3 6

**Space Complexity:** 4 bytes \* Execution length!



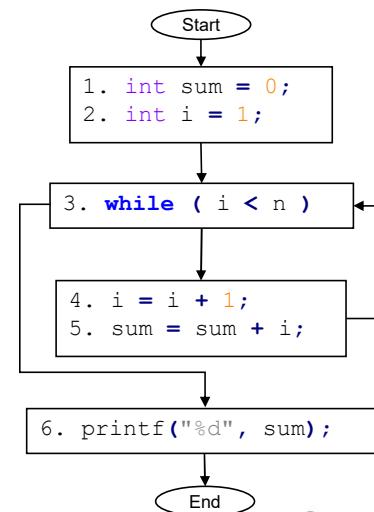
51

## Basic Block Level Tracing

```
void sumUp(int n) {
    int sum = 0;
    int i = 1;
    while ( i < n ) {
        i = i + 1;
        sum = sum + i;
    }
    printf("%d",
sum);
}
```

**Trace (n=6) :** 1 2 3 4 5 3 4 5 3 4 5 3 4 5 3 6

**BB Trace:** 1 3 4 3 4 3 4 3 4 3 6



52

## More Ideas

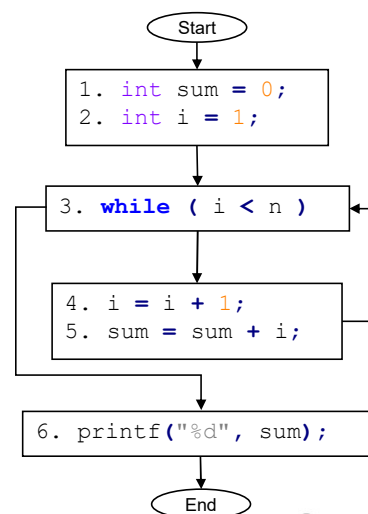
- Would function level tracing work?
  - Each trace entry is a function call & its parameters
  - No...
  - Function's behavior can be affected by global variables ☹
  - Cannot distinguish between nested functions or those called in sequence! ☹
- Predicate tracing

Fine-grained Trace	Predicate Trace
--------------------	-----------------

1 2 3 6	F
---------	---

1 2 3 4 5 3 6	T F
---------------	-----

- The trace is no longer randomly accessible
- Must start from the beginning to understand



53

## Lesson Summary

- Discussed how to employ the tracing process for dynamic analysis
- Discussed how to Apply Static and Dynamic BinaryInstrumentation
- Discussed how to Create a Pintool to utilize PIN
- Discussed Valgrind and QEMU techniques of instrumentation



54